

Student Performance Factors

Amy Fishencord

This dataset provides a comprehensive overview of various factors affecting student performance in exams. It includes information on study habits, attendance, parental involvement, and other aspects influencing academic success.

```
In [8]: #importing required libraries
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

#loading in data set
data = pd.read_csv('StudentPerformanceFactors.csv')
print(data.info())
data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 6607 entries, 0 to 6606
```

```
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	Hours_Studied	6607 non-null	int64
1	Attendance	6607 non-null	int64
2	Parental_Involvement	6607 non-null	object
3	Access_to_Resources	6607 non-null	object
4	Extracurricular_Activities	6607 non-null	object
5	Sleep_Hours	6607 non-null	int64
6	Previous_Scores	6607 non-null	int64
7	Motivation_Level	6607 non-null	object
8	Internet_Access	6607 non-null	object
9	Tutoring_Sessions	6607 non-null	int64
10	Family_Income	6607 non-null	object
11	Teacher_Quality	6529 non-null	object
12	School_Type	6607 non-null	object
13	Peer_Influence	6607 non-null	object
14	Physical_Activity	6607 non-null	int64
15	Learning_Disabilities	6607 non-null	object
16	Parental_Education_Level	6517 non-null	object
17	Distance_from_Home	6540 non-null	object
18	Gender	6607 non-null	object
19	Exam_Score	6607 non-null	int64

```
dtypes: int64(7), object(13)
```

```
memory usage: 1.0+ MB
```


```
None
```

Out [8]:

	Hours_Studied	Attendance	Parental_Involvement	Access_to_Resources	Extracurricular
0	23	84	Low	High	
1	19	64	Low	Medium	
2	24	98	Medium	Medium	
3	29	89	Low	Medium	
4	19	92	Medium	Medium	

Overview of Variables

In [9]: `#distribution plot`
`sns.displot(data['Exam_Score'], kde=True, color='pink', bins=10, aspect=1.5)`
`plt.title('Distribution of Exam Scores') #adding plot labels`
`plt.xlabel('Exam Scores')`
`plt.ylabel('Count')`
`plt.show()`



Most students score in the 60's and 70's

In [6]: `print(data.describe())`

	Hours_Studied	Attendance	Sleep_Hours	Previous_Scores	\
count	6607.000000	6607.000000	6607.000000	6607.000000	
mean	19.975329	79.977448	7.02906	75.070531	
std	5.990594	11.547475	1.46812	14.399784	
min	1.000000	60.000000	4.00000	50.000000	
25%	16.000000	70.000000	6.00000	63.000000	
50%	20.000000	80.000000	7.00000	75.000000	
75%	24.000000	90.000000	8.00000	88.000000	
max	44.000000	100.000000	10.00000	100.000000	

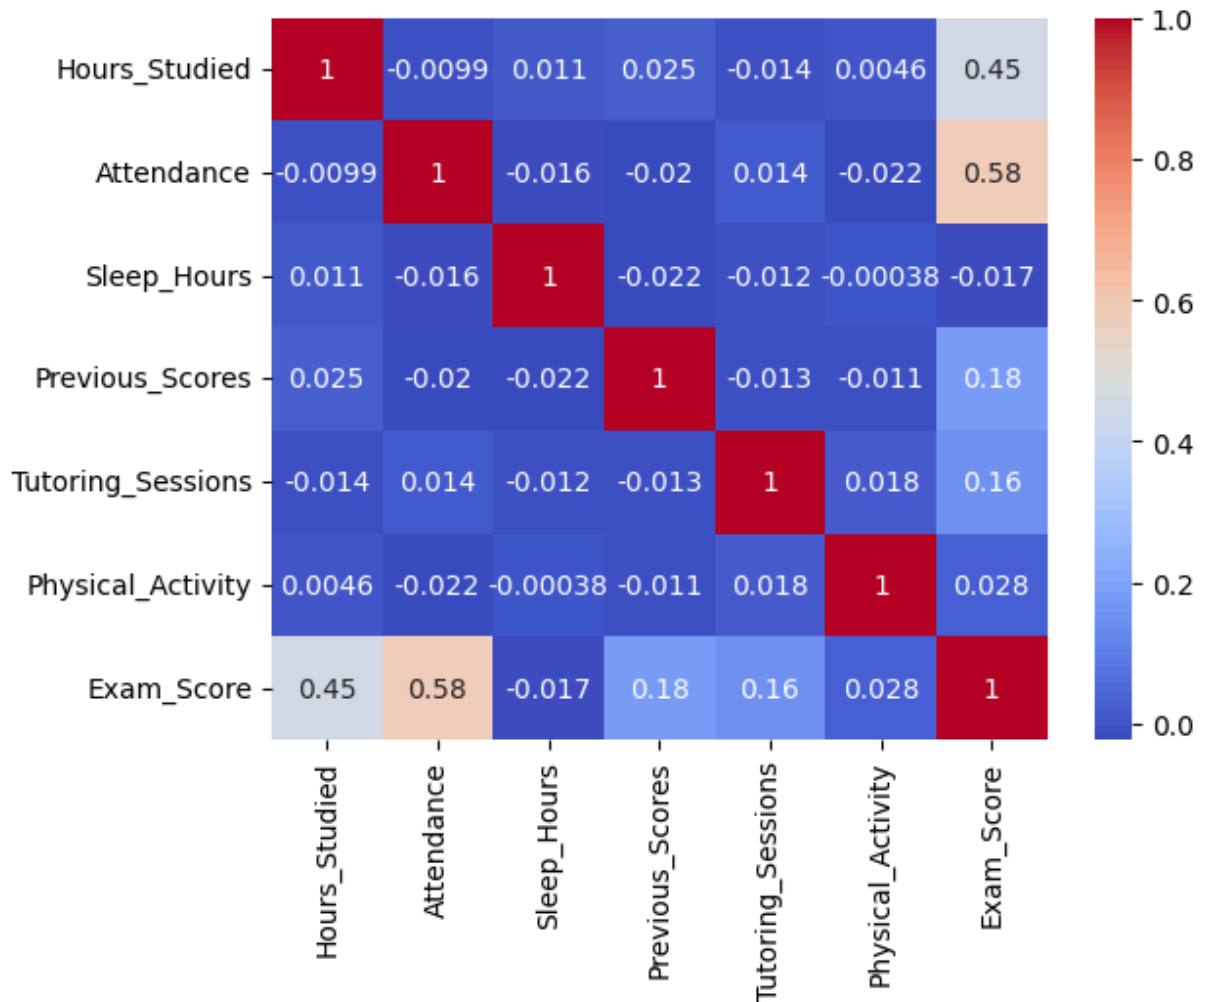
	Tutoring_Sessions	Physical_Activity	Exam_Score
count	6607.000000	6607.000000	6607.000000
mean	1.493719	2.967610	67.235659
std	1.230570	1.031231	3.890456
min	0.000000	0.000000	55.000000
25%	1.000000	2.000000	65.000000
50%	1.000000	3.000000	67.000000
75%	2.000000	4.000000	69.000000
max	8.000000	6.000000	101.000000

Question 1:

What factors have the most significant impact on students' exam scores?

```
In [12]: #correlation matrix
numeric_data = data.select_dtypes(include=[np.number])
corr_matrix = numeric_data.corr()


#add heatmap
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm")
plt.show()
```



Positive correlation between exam scores and other variables that we will be looking at further.

```
In [10]: #linear regression plot
sns.regplot(x='Exam_Score', y='Hours_Studied', data=data, color='deeppink',
plt.title('Scatter Plot of Exam Scores vs Hours Studied')
plt.xlabel('Exam Score')
plt.ylabel('Hours Studied')
plt.show()

#Calculated correlation coeff
correlation = data['Exam_Score'].corr(data['Hours_Studied'])
print(correlation)
```



0.4454549540752825

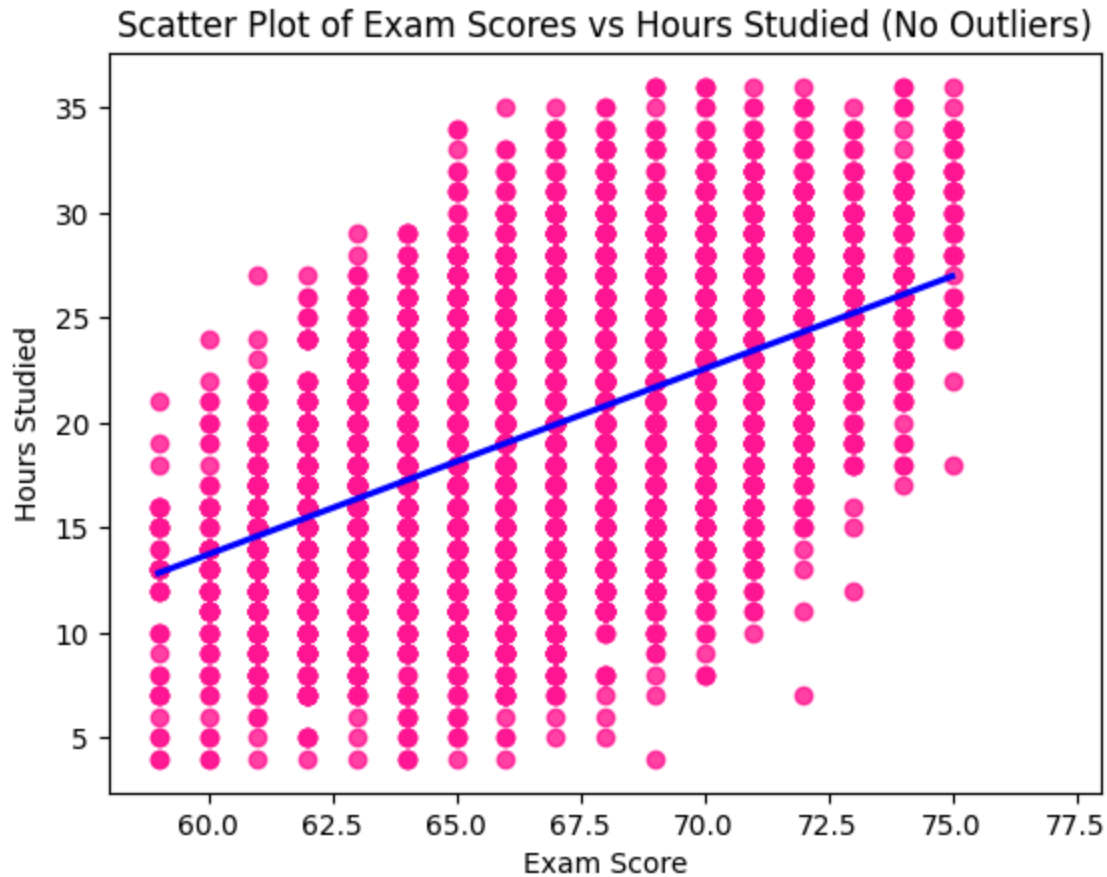
A calculated correlation coefficient of 0.45 showing a moderately positive correlation between exam scores and hours studied. Would this correlation be stronger without the outliers?

```
In [13]: #calculate IQR to remove outliers
Q1 = data[['Exam_Score', 'Hours_Studied']].quantile(0.25)
Q3 = data[['Exam_Score', 'Hours_Studied']].quantile(0.75)
IQR = Q3 - Q1

#filter out the outliers using the IQR method
filtered_data = data[~((data[['Exam_Score', 'Hours_Studied']] < (Q1 - 1.5 *
    (data[['Exam_Score', 'Hours_Studied']] > (Q3 + 1.5 *

#linear reg plot (without outliers)
sns.regplot(x='Exam_Score', y='Hours_Studied', data=filtered_data, color='de
plt.title('Scatter Plot of Exam Scores vs Hours Studied (No Outliers)')
plt.xlabel('Exam Score')
plt.ylabel('Hours Studied')
plt.xlim(58,78)
plt.show()

#calculated correlation coefficient (without outliers)
correlation = filtered_data['Exam_Score'].corr(filtered_data['Hours_Studied']
print("Correlation coefficient (without outliers):", correlation)
```




Correlation coefficient (without outliers): 0.494991169076997

When we remove our outliers in the exam score variable, our calculated correlation coefficient is 0.5. This shows a moderately positive correlation that is higher without the outliers.

```
In [14]: #linear regression plot
sns.regplot(x='Exam_Score', y='Attendance', data=data, color='deeppink', ci=
plt.title('Scatter Plot of Exam Scores vs Attendance')
plt.xlabel('Exam Score')
plt.ylabel('Attendance (% of Classes Attended)')
plt.ylim(10,110)
plt.show()

#correlation coefficient
correlation = data['Exam_Score'].corr(data['Attendance'])
print(correlation)
```



0.5810718633120647

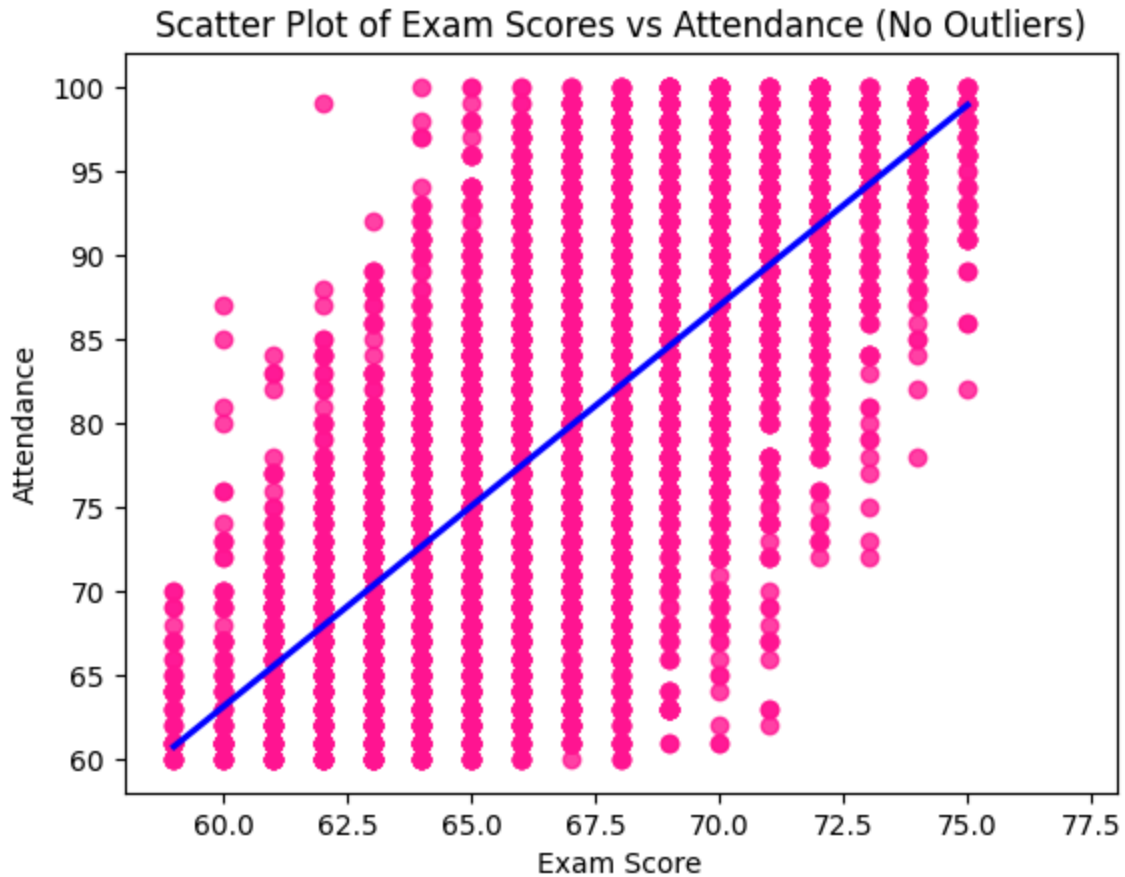
A calculated correlation coefficient of 0.58 showing a moderately positive correlation between exam scores and attendance.

```
In [15]: #calculate IQR to remove outliers
Q1 = data[['Exam_Score', 'Attendance']].quantile(0.25)
Q3 = data[['Exam_Score', 'Attendance']].quantile(0.75)
IQR = Q3 - Q1

#filter out the outliers using the IQR method
filtered_data = data[~((data[['Exam_Score', 'Attendance']] < (Q1 - 1.5 * IQR)
                       & (data[['Exam_Score', 'Attendance']] > (Q3 + 1.5 * IQR)))]

#linear reg plot (without outliers)
sns.regplot(x='Exam_Score', y='Attendance', data=filtered_data, color='deepskyblue')
plt.title('Scatter Plot of Exam Scores vs Attendance (No Outliers)')
plt.xlabel('Exam Score')
plt.ylabel('Attendance')
plt.xlim(58,78)
plt.show()

#calculated correlation coefficient (without outliers)
correlation = filtered_data['Exam_Score'].corr(filtered_data['Attendance'])
print("Correlation coefficient (without outliers):", correlation)
```




Correlation coefficient (without outliers): 0.6738707423757742

When we remove our outliers in the attendance variable, our calculated correlation coefficient is 0.67. This shows a high correlation that is stronger without the outliers.

```
In [16]: #linear regression plot
sns.regplot(x='Exam_Score', y='Sleep_Hours', data=data, color='deeppink', ci=
plt.title('Scatter Plot of Exam Scores vs Sleep Hours')
plt.xlabel('Exam Score')
plt.ylabel('Sleep Hours')
plt.show()

#correlation coefficient
correlation = data['Exam_Score'].corr(data['Sleep_Hours'])
print(correlation)
```





-0.01702162857150254

A calculated correlation coefficient shows a weak negative correlation close to zero. The amount of hours of sleep you get will have almost no impact on your exam score

```
In [17]: #linear regression
sns.regplot(x='Exam_Score', y='Tutoring_Sessions', data=data, color='deeppir')
plt.title('Scatter Plot of Exam Scores vs Tutoring Sessions')
plt.xlabel('Exam Score')
plt.ylabel('Tutoring Sessions')
plt.show()

#correlation coefficient
correlation = data['Exam_Score'].corr(data['Tutoring_Sessions'])
print(correlation)
```




0.15652518539225332

A calculated correlation coefficient of 0.15 shows a weak positive correlation. The amount of tutoring sessions you attend will not have much of an impact on your exam score.

```
In [18]: #Linear Reg plot
sns.regplot(x='Exam_Score', y='Physical_Activity', data=data, color='deeppir')
plt.title('Scatter Plot of Exam Scores vs Physical Activity')
plt.xlabel('Exam Score')
plt.ylabel('Physical Activity')
plt.show()

#Correlation coefficient
correlation = data['Exam_Score'].corr(data['Physical_Activity'])
print(correlation)
```



0.02782443618025749

A calculated correlation coefficient of 0.29 shows a low correlation. It is not likely that changes in physical activity do not significantly predict changes in exam scores.

To answer our question the factors that impact students exam scores the most are hours studied and attendance. Attendance affects students exam scores the most.



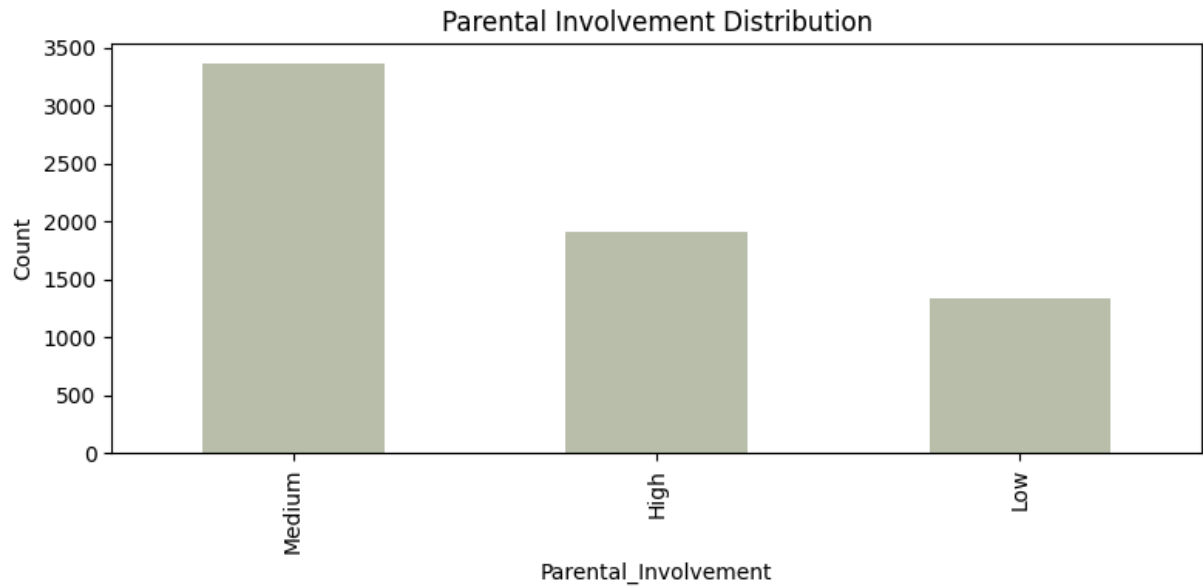
Question 2: How do family related factors influence exam scores?

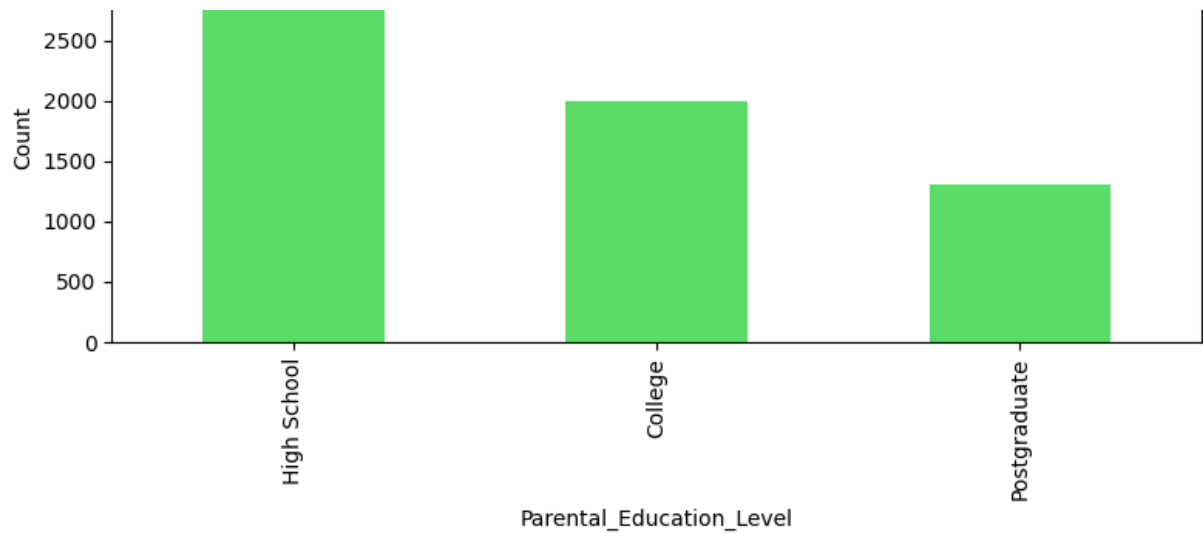
```
In [19]: categorical_vars = ['Parental_Involvement', 'Access_to_Resources', 'Family_I

#create subplots for all variables
fig, axes = plt.subplots(len(categorical_vars), 1, figsize=(8, len(categoric

#loop through each variable and create a bar plot
for i, var in enumerate(categorical_vars):
    counts = data[var].value_counts() #get counts of each category
    counts.plot.bar(ax=axes[i], color=np.random.rand(3,), alpha=0.7)
    axes[i].set_title(f'{var.replace("_", " ").title()} Distribution') #addi
    axes[i].set_ylabel('Count')

plt.tight_layout() #adjust layout
plt.show()
```





```
In [20]: fig, axes = plt.subplots(len(categorical_vars), 1, figsize=(9, 17))

#loop through each variable and create a boxplot
for i, var in enumerate(categorical_vars):
    sns.boxplot(x=data[var], y=data['Exam_Score'], ax=axes[i], palette='Set2')
    axes[i].set_title(f'{var.replace("_", " ").title()} vs Exam Score') #title
    axes[i].set_xlabel(f'{var.replace("_", " ").title()}') #xaxis label
    axes[i].set_ylabel('Exam Score') #yaxis label

#adjust layout and show the plot
plt.tight_layout()
plt.show()
```

```
/tmp/ipykernel_632366/1621321702.py:5: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=data[var], y=data['Exam_Score'], ax=axes[i], palette='Set2')  
/tmp/ipykernel_632366/1621321702.py:5: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.



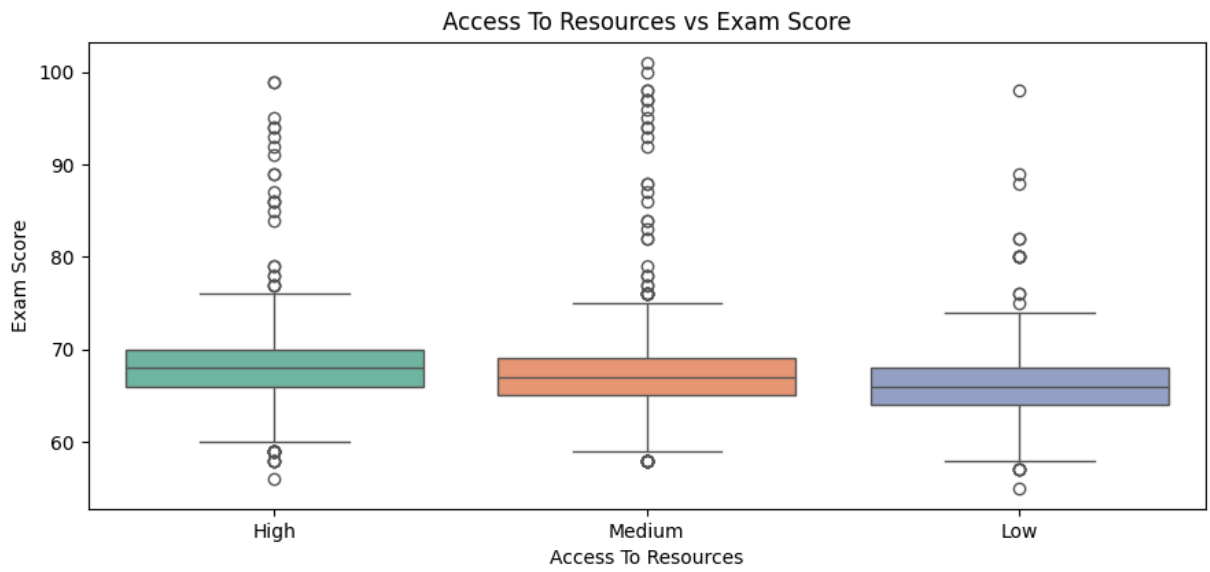
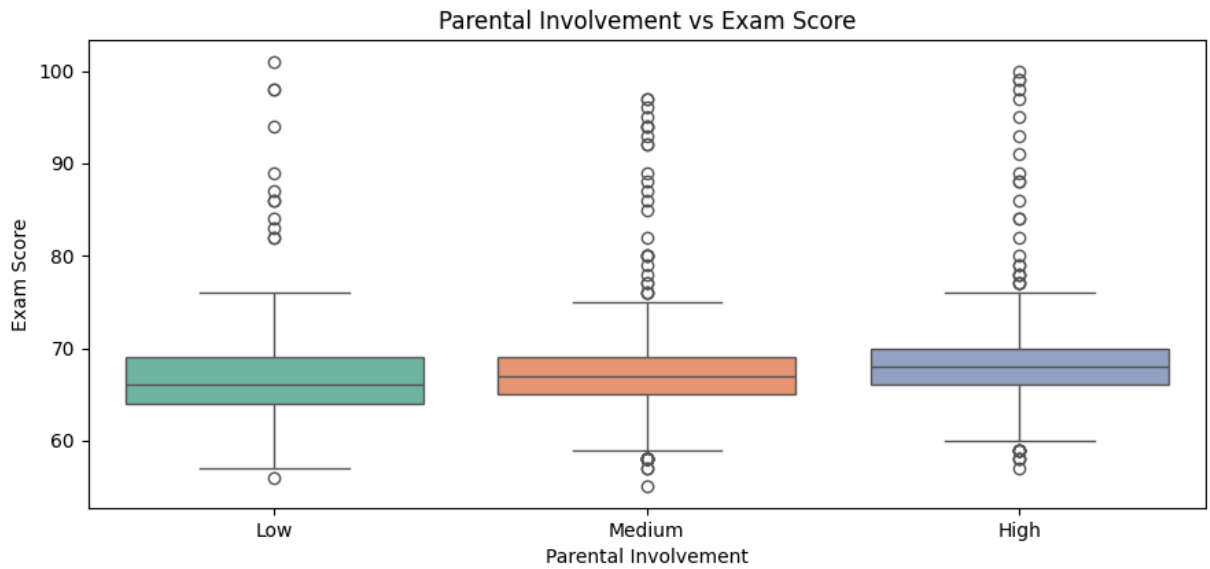
```
sns.boxplot(x=data[var], y=data['Exam_Score'], ax=axes[i], palette='Set2')  
/tmp/ipykernel_632366/1621321702.py:5: FutureWarning:
```


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=data[var], y=data['Exam_Score'], ax=axes[i], palette='Set2')  
/tmp/ipykernel_632366/1621321702.py:5: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=data[var], y=data['Exam_Score'], ax=axes[i], palette='Set2')
```



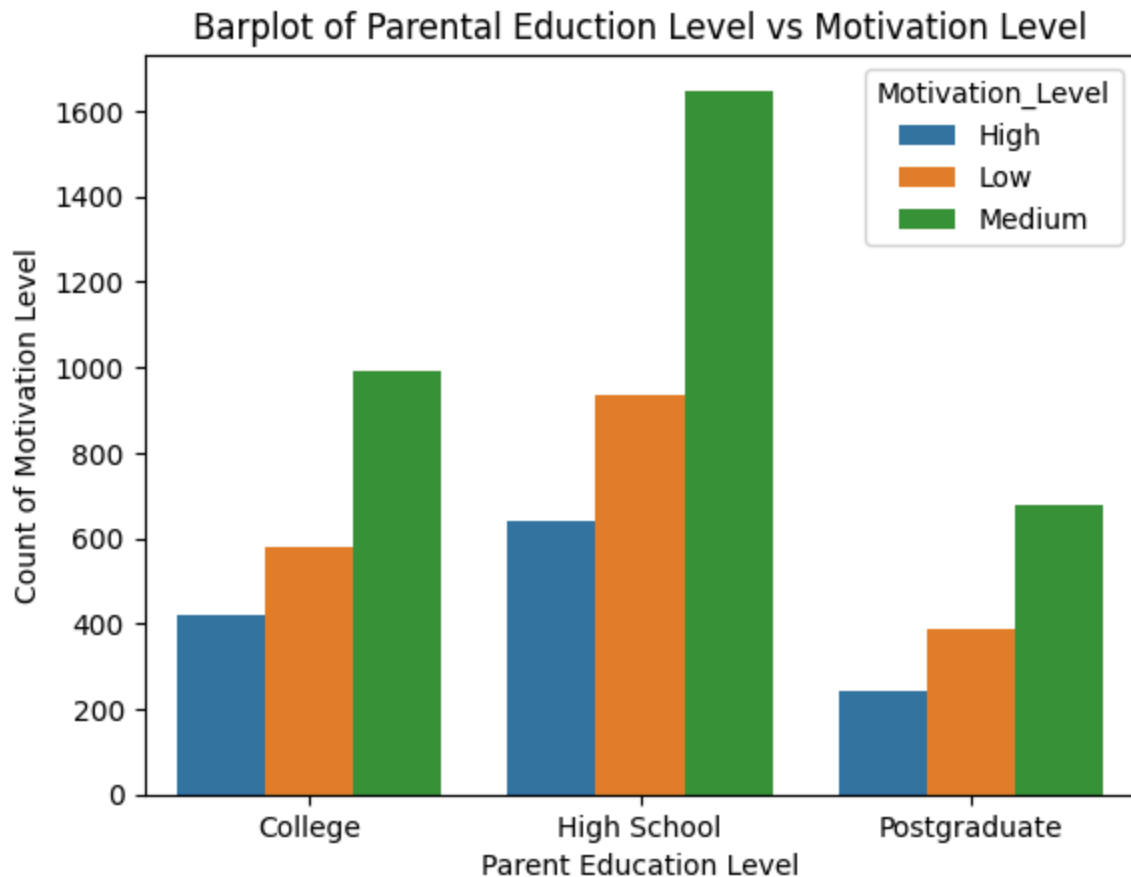


To answer our question, almost all students who have higher parent involvement, access to resources, family income, and education level tend to have higher exam scores. We see this by focusing on the mean values, which show a positive trend, and by examining our outlier variables, which further emphasize the stronger performance of students with these factors. The presence of outliers in the data may indicate extreme cases, but the overall pattern remains consistent with the influence of these variables on exam scores.

Question 3: Do parents education levels affect students motivational levels?

```
In [21]: #calculate the counts of Motivation_Level by Parental_Education_Level
count_data = data.groupby(['Parental_Education_Level', 'Motivation_Level']).

#create barplot
sns.barplot(x='Parental_Education_Level', y='count', hue='Motivation_Level',
plt.title('Barplot of Parental Education Level vs Motivation Level')
plt.xlabel('Parent Education Level')
plt.ylabel('Count of Motivation Level')
plt.show()
```

Parents education levels do not have an affect on students motivational levels, most student have an average stress level.

Question 4: Do parents education levels affect the type of school students go to?

```
In [22]: #calculate the counts of School_Type by Parental_Education_Level
count_data = data.groupby(['Parental_Education_Level', 'School_Type']).size()

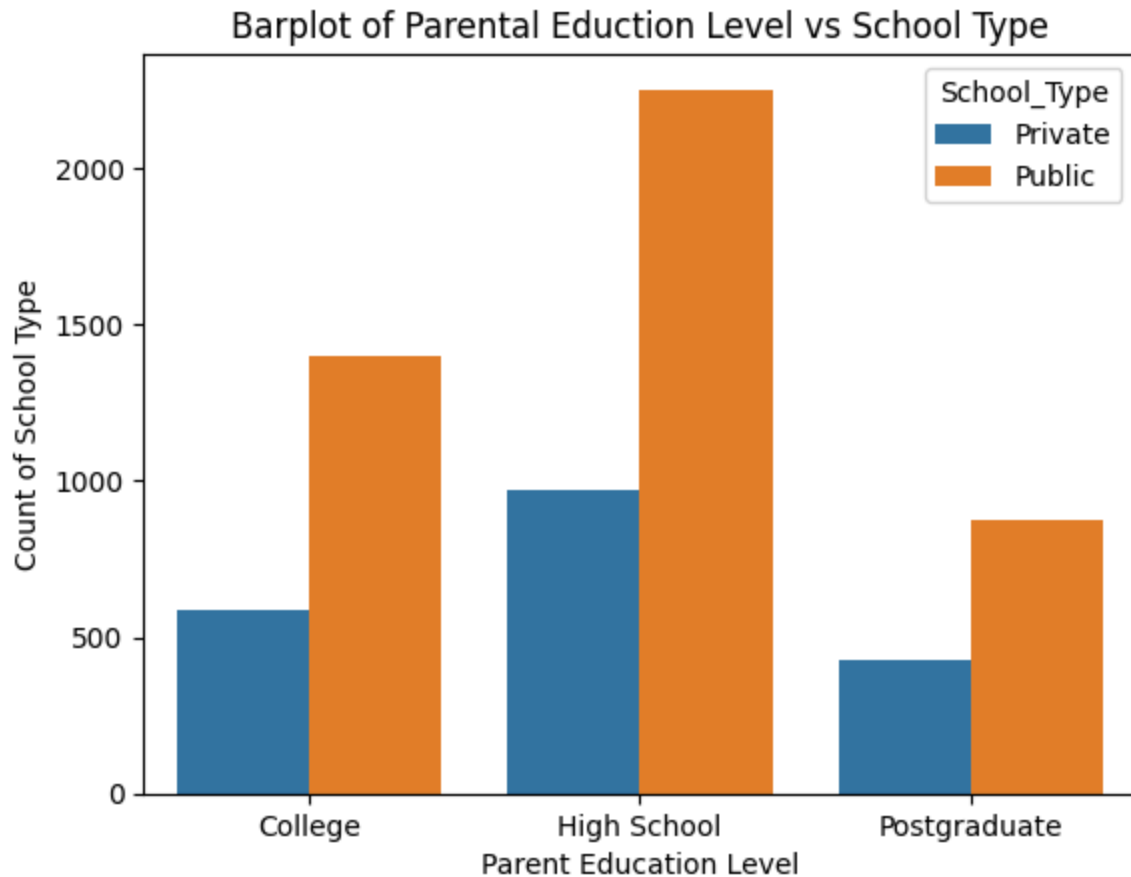
#pivot the table to restructure counts
pivot_table = count_data.pivot(index='Parental_Education_Level', columns='School_Type')

#add a column for the public to private ratio
pivot_table['Public_to_Private_Ratio'] = pivot_table['Public'] / pivot_table['Private']

print(pivot_table)

#create barplot
sns.barplot(x='Parental_Education_Level', y='count', hue='School_Type', data=count_data)
plt.title('Barplot of Parental Education Level vs School Type')
plt.xlabel('Parent Education Level')
plt.ylabel('Count of School Type')
plt.show()
```

School_Type	Private	Public	Public_to_Private_Ratio
Parental_Education_Level			
College	589	1400	2.376910
High School	972	2251	2.315844
Postgraduate	428	877	2.049065



Parent educational levels do not affect the type of school students go to. The ratio of public to private school are similar by all education levels and most students attend public schools.

Question 5: Does family income affect the type of school students go to?

```
In [23]: #calculate the counts of School_Type by Famil
count_data = data.groupby(['Family_Income', 'School_Type']).size().reset_in

#pivot the table to restructure counts
pivot_table = count_data.pivot(index='Family_Income', columns='School_Type',

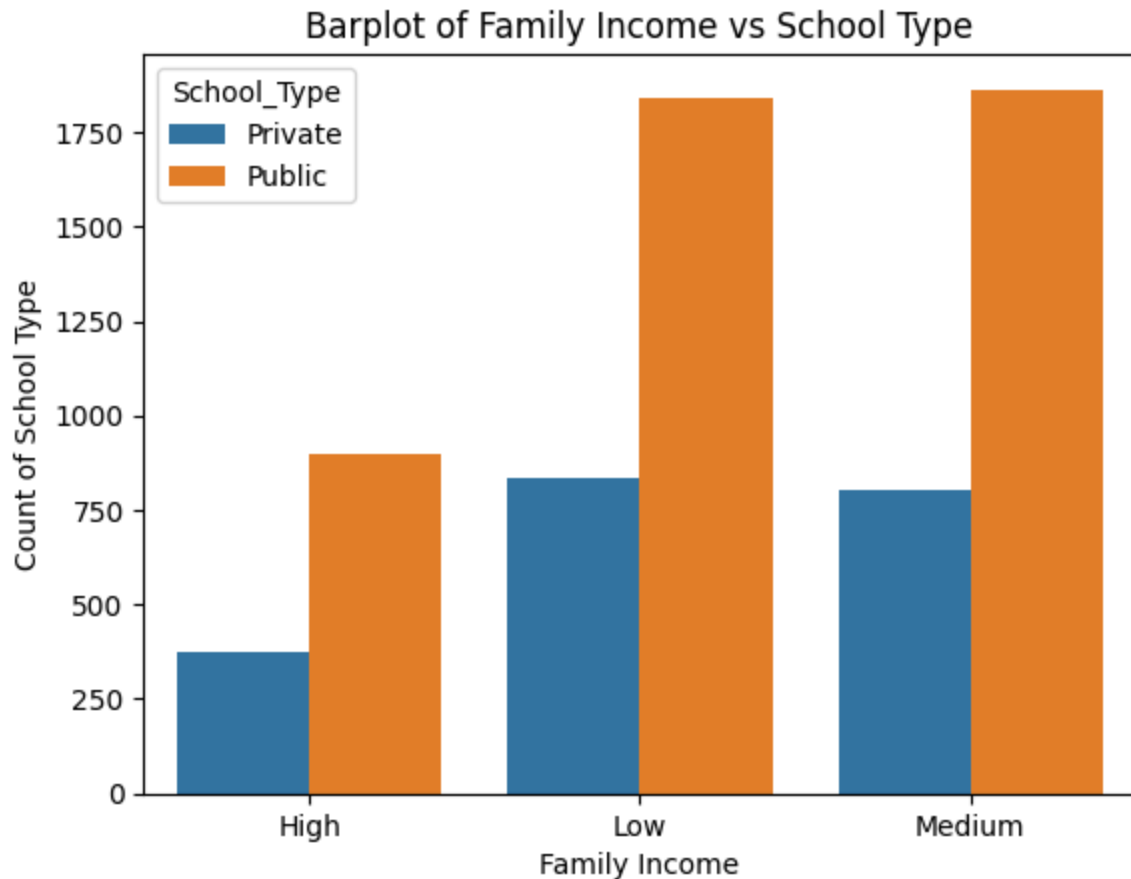
#add a column for the public to private ratio
pivot_table['Public_to_Private_Ratio'] = pivot_table['Public'] / pivot_table

print(pivot_table)

#create barplot
```

```
sns.barplot(x='Family_Income', y='count', hue='School_Type', data=count_data)
plt.title('Barplot of Family Income vs School Type')
plt.xlabel('Family Income')
plt.ylabel('Count of School Type')
plt.show()
```

School_Type	Private	Public	Public_to_Private_Ratio
Family_Income			
High	373	896	2.402145
Low	833	1839	2.207683
Medium	803	1863	2.320050



Family income does not affect the type of school students go to. The ratios of public to private schools are also similar by family income and most students attend public school.

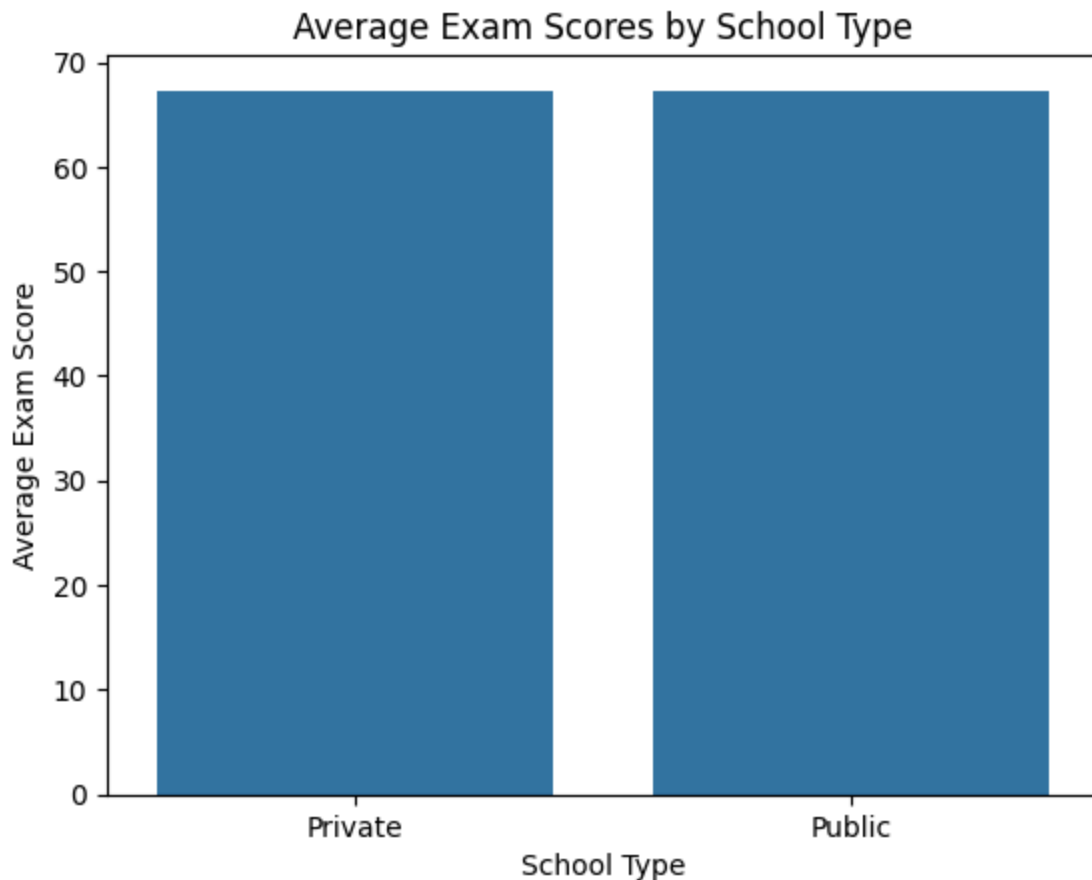
Question 6: Which type of schools have the higher exam scores on average?

```
In [24]: #group by school type and calc mean exam score
average_scores = data.groupby('School_Type')['Exam_Score'].mean().reset_index()

#create barplot
sns.barplot(x='School_Type', y='Exam_Score', data=average_scores)

plt.title('Average Exam Scores by School Type') #adding title and labels
```

```
plt.xlabel('School Type')
plt.ylabel('Average Exam Score')
plt.show()
```



As we can see neither public or private schools are producing higher scores. Public and private schools have the same average exam scores, showing the type of school you go to will most likely not have an impact on your exam score.

Question 7: How do motivation levels and study time interact?

```
In [25]: plt.figure(figsize=(10, 6))

#scatterplot grouped by motivational level
sns.lmplot(x='Hours_Studied', y='Exam_Score', hue='Motivation_Level', data=c

#labels and title
plt.title('Interaction between Motivation Level and Study Time on Exam Score
plt.xlabel('Study Time (hours)')
plt.ylabel('Exam Scores')
plt.legend(title='Motivation Level')
plt.show()

#calculate correlation coefficients for each motivation level
for level in data['Motivation_Level'].unique():
```

```
subset = data[data['Motivation_Level'] == level]
correlation = subset['Hours_Studied'].corr(subset['Exam_Score'])
print(f"Correlation for Motivation Level {level}: {correlation:.2f}")
```

<Figure size 1000x600 with 0 Axes>



Correlation for Motivation Level Low: 0.47

Correlation for Motivation Level Medium: 0.44

Correlation for Motivation Level High: 0.42

Low Motivation Level (0.47): A moderate positive correlation. This suggests that for students with low motivation, there is a noticeable relationship between study time and exam scores—students who study more tend to perform better, though the relationship is not very strong.

Medium Motivation Level (0.44): Another moderate positive correlation, slightly lower than for low motivation. This implies a similar trend, but with slightly weaker predictive power.

High Motivation Level (0.42): The weakest positive correlation among the three levels. This suggests that for highly motivated students, study time is slightly less predictive of exam scores compared to those with medium or low motivation.

Predictive Model

Next, I will be creating a predictive model to predict the a students exam score based on their attendance and hours studied.

```
In [26]: #calculate IQR to remove outliers
Q1 = data[['Attendance', 'Hours_Studied', 'Exam_Score']].quantile(0.25)
```

```

Q3 = data[['Attendance', 'Hours_Studied', 'Exam_Score']].quantile(0.75)
IQR = Q3 - Q1

#filter out outliers using IQR method
filtered_data = data[~(
    ((data[['Attendance', 'Hours_Studied', 'Exam_Score']] < (Q1 - 1.5 * IQR)
    (data[['Attendance', 'Hours_Studied', 'Exam_Score']] > (Q3 + 1.5 * IQR)
))]

#set x (features) and y (target)
X = filtered_data[['Attendance', 'Hours_Studied']]
y = filtered_data[['Exam_Score']]

```

Next, I will split the data into training and test sets. Will use 80% of the data for training and 20% of the data for testing to evaluate the model's performance.

```

In [27]: from sklearn.model_selection import train_test_split

#split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

```

Next, I will use the LinearRegression class to fit the model.

```

In [28]: from sklearn.linear_model import LinearRegression

#initialize model
model = LinearRegression()

#train model
model.fit(X_train, y_train)

```

```

Out[28]: ▼ LinearRegression ⓘ ⓘ
LinearRegression()

```

Next, I will apply the model to make a prediction.

```

In [29]: #make predictions on the test set
y_pred = model.predict(X_test)

```

Next, I will test the performance of the model.

```

In [31]: from sklearn.metrics import mean_absolute_error, r2_score

mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Absolute Error (MAE):", mae)
print("R-squared (R2):", r2)

```

```

Mean Absolute Error (MAE): 1.3290443509773582
R-squared (R2): 0.7289459523708635


```

The Mean Absolute Error (MAE) represents the average difference between the predicted and actual exam scores. In this case, the MAE of 1.33 indicates that the model's predictions are, on average, 1.33 points away from the actual scores. This suggests the model provides fairly accurate predictions, especially when considering the scale of exam scores.

The R-squared value (0.73) indicates that 73% of the variation in exam scores is explained by the model's predictors (attendance and hours studied). This suggests that the model captures a significant portion of the relationship between these variables and exam performance.

Actual vs Predicted Comparison

```
In [33]: plt.scatter(y_test, y_pred, color = 'pink', alpha=0.8, label='Predicted vs Actual')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', linestyle='dashed', label='Perfect Prediction')
plt.title('Actual vs Predicted')
plt.xlabel('Actual Exam Scores')
plt.ylabel('Predicted Exam Scores')
plt.legend()
plt.grid(True)
plt.show()
```



The pink points on the graph represent the actual exam scores on the x axis vs the predicted exam score on the y axis. The red dashed line represents a perfect prediction where the predicted score exactly matches the actual score. If the model predicted perfectly, all points would lie exactly on this line. If a data point lies exactly on this line, it means the model predicted that score perfectly, points above the line mean the model predicted a lower score than the actual value (underprediction), and points below the line mean the model predicted a higher score than the actual value (overprediction). Most data points are scattered around the perfect fit line representing a good model.

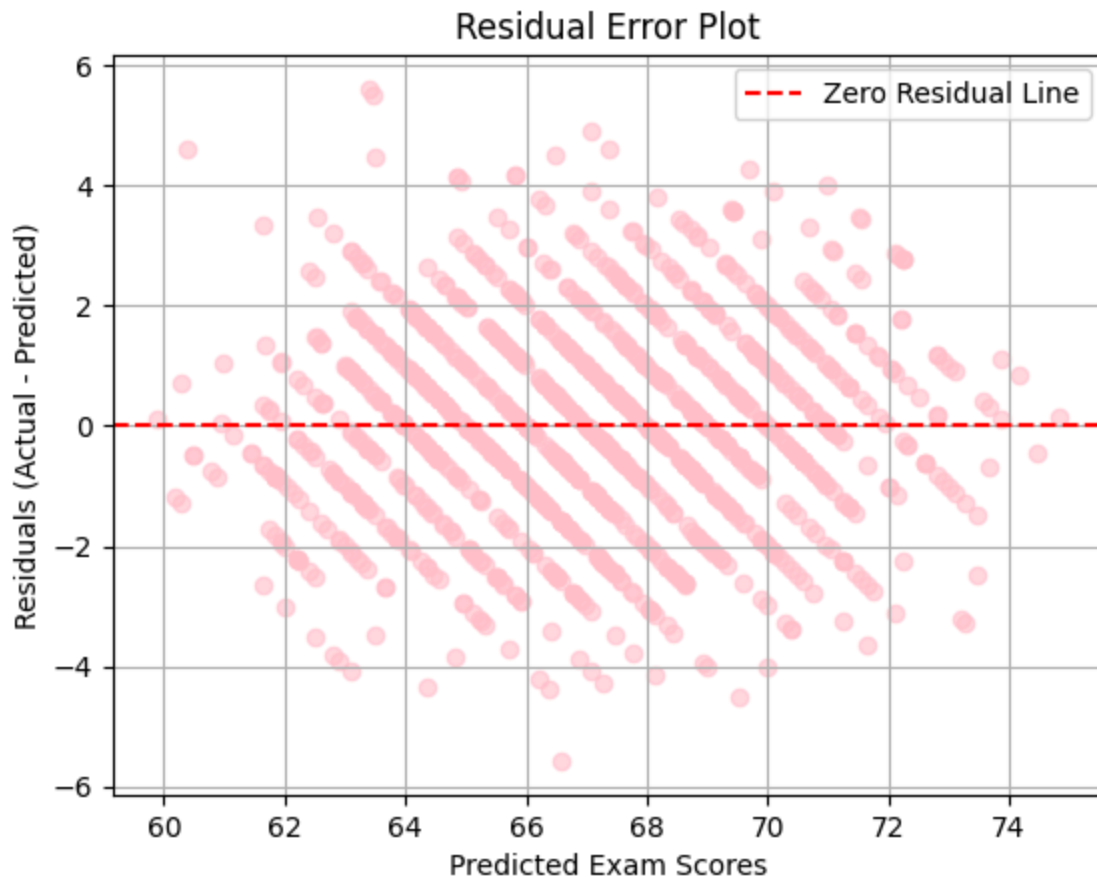
Residual Plot

Next, I will create a residual plot to show the difference between the actual values and the predicted values.

```
In [35]: #calculate residuals
residuals = y_test - y_pred

#plot residual error
plt.scatter(y_pred, residuals, color = 'pink', alpha = 0.6)
plt.axhline(y=0, color='red', linestyle='--', label='Zero Residual Line')
plt.title('Residual Error Plot')
plt.xlabel('Predicted Exam Scores')
plt.ylabel('Residuals (Actual - Predicted)')
plt.grid(True)
plt.legend()
```

```
Out[35]: <matplotlib.legend.Legend at 0x72d506f97cb0>
```

The red dashed line represents the zero residual line, which indicates where the predicted values would exactly match the actual values. Most of the points are scattered around this line with no clear pattern, indicating that the model's residuals are randomly distributed. This suggests that the model has effectively captured the relationship between the features (attendance and hours studied) and the exam scores, without any significant bias or unaccounted patterns.

To understand the model easier, let's provide a simple example using it. Let us say student 1 attends class 85 percent of the time and studies for 15 hours a week. What is their predicted exam score?

```
In [40]: student_data = pd.DataFrame([[85, 15]], columns=['Attendance', 'Hours_Studied'])
         predicted_score = model.predict(student_data)

         print(f"Predicted exam score for Student 1 (85% attendance, 15 hours studied): {predicted_score}")
```

Predicted exam score for Student 1 (85% attendance, 15 hours studied): 66.60

Thank you!!